

```

import os
import glob
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import re
from PIL import Image
from tqdm.notebook import tqdm

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms
from torchvision.models import resnet18

torch.random.manual_seed(42)
np.random.seed(42)

#Dataset Preparation

class EuropaDataset(Dataset):
    def __init__(self, image_path, file_path, file_list):
        self.image_path= image_path
        self.file_path= file_path
        self.images_dict= self.image_process()
        self.labels_dict= self.text_process()
        self.file_list= file_list

    def __len__(self):
        return len(self.file_list)

    def __getitem__(self, index):
        filename= self.file_list[index]
        return self.images_dict[filename], self.labels_dict[filename]

    def image_process(self):
        transform= transforms.Compose(
            [
                transforms.Resize((50, 200)),
                transforms.ToTensor(),
                transforms.Normalize(mean=(0.485, 0.456, 0.406),
std=(0.229, 0.224, 0.225))
            ]
        )
        images=
{val:transform(Image.open(os.path.join(self.image_path,
val))).convert("RGB"))
    for val in os.listdir(self.image_path)}

```

```

        return images

    def text_process(self):
        with open(self.file_path) as f:
            text= f.read().split("\n")

            labels_dict={}
            for i in range(len(text)):
                ind= text[i].find("=")
                filename= text[i][:ind].strip()
                cont= text[i][ind+1:].strip()

                cont_cleaned = re.sub(r'[^a-zA-Z0-9 ]+', '', cont)
                labels_dict[filename]= cont_cleaned
            return labels_dict

with open("./data_words.txt") as f:
    text= f.read().split("\n")

file_list=[]
for i in range(len(text)):
    try:
        ind= text[i].index("=")
        filename= text[i][:ind].strip()
        file_list.append(filename)
    except:
        continue
train_size= int(len(file_list)*0.8)
train_file_inds= np.random.choice(len(file_list), size=(train_size, ),
replace=False)
train_file_names= [file_list[i] for i in train_file_inds]
test_file_names= [val for val in file_list if val not in
train_file_names]

train_set= EuropaDataset("./cropped_words", "./data_words.txt",
train_file_names)
test_set= EuropaDataset("./cropped_words", "./data_words.txt",
test_file_names)
train_loader= DataLoader(train_set, 16)
test_loader= DataLoader(test_set, 16)

#Utility Functions
criterion = nn.CTCLoss(blank=0)
alphabet= [val for val in "-
0123456789ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz"]
encoder_map= {ind:val for val, ind in enumerate(alphabet)}
decoder_map= {val:ind for val, ind in enumerate(alphabet)}

def encode_text(text_batch):
    text_lengths=[]

```

```

    for text in text_batch:
        text_lengths.append(len(text))
    strv= "".join(text_batch)
    encoding=[encoder_map[char] for char in strv]
    text_lengths= torch.IntTensor(text_lengths)
    encoding= torch.IntTensor(encoding)
    return encoding, text_lengths

def calculate_CTCLoss(output_logits, text_labels, device):
    log_probs= nn.functional.log_softmax(output_logits, 2)
    input_lengths= torch.full(size=(log_probs.size(1),),
    fill_value=log_probs.size(0), dtype=torch.int32).to(device)
    targets, target_lengths= encode_text(text_labels)
    return criterion(log_probs, targets, input_lengths,
    target_lengths)

def decode_predictions(text_batch_logits):
    text_batch_tokens = F.softmax(text_batch_logits, 2).argmax(2) #
    [T, batch_size]
    text_batch_tokens = text_batch_tokens.numpy().T # [batch_size, T]

    text_batch_tokens_new = []
    for text_tokens in text_batch_tokens:
        text = [decoder_map[idx] for idx in text_tokens]
        text = "".join(text)
        text_batch_tokens_new.append(text)

    return text_batch_tokens_new

def remove_duplicates(text):
    if len(text) > 1:
        letters = [text[0]] + [letter for idx, letter in
    enumerate(text[1:], start=1) if text[idx] != text[idx-1]]
    elif len(text) == 1:
        letters = [text[0]]
    else:
        return ""
    return "".join(letters)

def correct_prediction(word):
    parts = word.split("-")
    parts = [remove_duplicates(part) for part in parts]
    corrected_word = "".join(parts)
    return corrected_word

batch_size = 16
num_chars = len(encoder_map)
rnn_hidden_size = 256
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(device)

```

```
cuda
```

```
resnet = resnet18(pretrained=True)
```

```
c:\Users\deban\AppData\Local\Programs\Python\Python312\Lib\site-packages\torchvision\models\_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
```

```
warnings.warn(
```

```
c:\Users\deban\AppData\Local\Programs\Python\Python312\Lib\site-packages\torchvision\models\_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use `weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
```

```
warnings.warn(msg)
```

```
class CRNN(nn.Module):
```

```
    def __init__(self, num_chars, rnn_hidden_size=256):
```

```
        super(CRNN, self).__init__()
```

```
        self.num_chars = num_chars
```

```
        self.rnn_hidden_size = rnn_hidden_size
```

```
        self.seq = nn.Sequential(*list(resnet.children())[:7]).to("cpu")
```

```
        self.linear1 = nn.Sequential(nn.Linear(1024, 512),
                                     nn.Dropout(0.2),
                                     nn.ReLU(),
                                     nn.Linear(512, 256))
```

```
        self.rnn1 = nn.LSTM(input_size=rnn_hidden_size,
                             hidden_size=rnn_hidden_size,
                             bidirectional=True,
                             batch_first=True,
                             num_layers=2,
                             dropout=0.2)
```

```
        self.linear2 = nn.Linear(self.rnn_hidden_size*2, num_chars)
```

```
    def forward(self, X):
```

```
        X= self.seq(X)
```

```
        X= X.permute(0,3,1,2)
```

```
        b, t, _, _ = X.shape
```

```
        X= X.view(b, t, -1)
```

```
        X= self.linear1(X)
```

```
        X, _ = self.rnn1(X)
```

```

        X = self.linear2(X)
        X = X.permute(1, 0, 2)
        return X

def weights_init(m):
    classname = m.__class__.__name__
    if type(m) in [nn.Linear, nn.Conv2d, nn.Conv1d]:
        torch.nn.init.xavier_uniform_(m.weight)
        if m.bias is not None:
            m.bias.data.fill_(0.01)
    elif classname.find('BatchNorm') != -1:
        m.weight.data.normal_(1.0, 0.02)
        m.bias.data.fill_(0)

num_epochs = 50
lr = 0.001
weight_decay = 1e-3
clip_norm = 5

crnn = CRNN(num_chars, rnn_hidden_size=rnn_hidden_size)
crnn.apply(weights_init)
crnn = crnn.to(device)

optimizer = optim.Adam(crnn.parameters(), lr=lr,
weight_decay=weight_decay)
lr_scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer,
patience=5)

epoch_losses = []
test_epoch_losses = []

for epoch in range(1, num_epochs+1):
    crnn.train()
    epoch_loss_list = []
    for image_batch, text_batch in train_loader:
        optimizer.zero_grad()
        text_batch_logits = crnn(image_batch.to(device))
        loss = calculate_CTCLoss(text_batch_logits, text_batch,
device)
        iteration_loss = loss.item()

        if np.isnan(iteration_loss) or np.isinf(iteration_loss):
            continue

        epoch_loss_list.append(iteration_loss)
        loss.backward()
        nn.utils.clip_grad_norm_(crnn.parameters(), clip_norm)
        optimizer.step()
    epoch_loss = np.mean(epoch_loss_list)
    epoch_losses.append(epoch_loss)

```

```

lr_scheduler.step(epoch_loss)

crnn.eval()
test_epoch_loss_list = []
with torch.inference_mode():
    for image_batch, text_batch in test_loader:
        text_batch_logits = crnn(image_batch.to(device))
        loss = calculate_CTCLoss(text_batch_logits, text_batch,
device)

        iteration_loss = loss.item()
        if np.isnan(iteration_loss) or np.isinf(iteration_loss):
            continue
        test_epoch_loss_list.append(iteration_loss)
test_epoch_loss = np.mean(test_epoch_loss_list)
test_epoch_losses.append(test_epoch_loss)


print("Epoch:{} Train Loss:{} Test Loss: {}".format(epoch,
epoch_loss, test_epoch_loss))

```

Epoch:1	Train Loss:3.4731808347565134	Test Loss: 3.0107797839926937
Epoch:2	Train Loss:2.458435390098243	Test Loss: 1.8751136368305987
Epoch:3	Train Loss:1.5842683704846214	Test Loss: 1.3607881101402077
Epoch:4	Train Loss:1.2606917689861863	Test Loss: 1.1550536779363183
Epoch:5	Train Loss:1.0982734324259051	Test Loss: 1.0114679811774074
Epoch:6	Train Loss:0.9824755316430872	Test Loss: 0.9863931872957462
Epoch:7	Train Loss:0.9242007535111391	Test Loss: 0.9218032123614462
Epoch:8	Train Loss:0.859545601719019	Test Loss: 0.900417598863366
Epoch:9	Train Loss:0.8191549528158453	Test Loss: 0.8487451852058352
Epoch:10	Train Loss:0.7852462433456804	Test Loss: 0.8374804900419758
Epoch:11	Train Loss:0.7655119483835959	Test Loss: 0.8355328032034934
Epoch:12	Train Loss:0.7431564062334705	Test Loss: 0.8175519360481082
Epoch:13	Train Loss:0.7242213743297677	Test Loss: 0.814755421721567
Epoch:14	Train Loss:0.7066529675819087	Test Loss: 0.816472921218421
Epoch:15	Train Loss:0.6918471796375713	Test Loss: 0.7859431503764911
Epoch:16	Train Loss:0.6724498062612908	Test Loss: 0.789270181208849
Epoch:17	Train Loss:0.6634302640217914	Test Loss: 0.7826848493522198
Epoch:18	Train Loss:0.6546809032202908	Test Loss: 0.8038327762332087
Epoch:19	Train Loss:0.6378656679292044	Test Loss: 0.7971760447219762
Epoch:20	Train Loss:0.6437949247266117	Test Loss: 0.8094912817176705
Epoch:21	Train Loss:0.6322848795561129	Test Loss: 0.8097754651716547
Epoch:22	Train Loss:0.6239831251818597	Test Loss: 0.7985463232578688
Epoch:23	Train Loss:0.6188764363883785	Test Loss: 0.8120800104409341
Epoch:24	Train Loss:0.6134428479526032	Test Loss: 0.7951186394757631
Epoch:25	Train Loss:0.6111450459683341	Test Loss: 0.8266358675298543
Epoch:26	Train Loss:0.605249994850615	Test Loss: 0.792989252394058
Epoch:27	Train Loss:0.6049732050328164	Test Loss: 0.7809182388355603
Epoch:28	Train Loss:0.5983000218511768	Test Loss: 0.7750753135956162
Epoch:29	Train Loss:0.5909290796903331	Test Loss: 0.7950089076195904
Epoch:30	Train Loss:0.5846650014867623	Test Loss: 0.8050797673150839
Epoch:31	Train Loss:0.5842842684051637	Test Loss: 0.7936258008523798

Epoch:32	Train Loss:0.5800906461641264	Test Loss: 0.7940571243410866
Epoch:33	Train Loss:0.5823490197364793	Test Loss: 0.7913100133761476
Epoch:34	Train Loss:0.5711916005176505	Test Loss: 0.7848888492089441
Epoch:35	Train Loss:0.5742328241538773	Test Loss: 0.7712164230529636
Epoch:36	Train Loss:0.5682428265088483	Test Loss: 0.7966168045163384
Epoch:37	Train Loss:0.5661663965103729	Test Loss: 0.802892906504996
Epoch:38	Train Loss:0.5584302660701663	Test Loss: 0.7946912133670682
Epoch:39	Train Loss:0.5587257727980613	Test Loss: 0.7757290671201976
Epoch:40	Train Loss:0.551216513617187	Test Loss: 0.7642784455942141
Epoch:41	Train Loss:0.555831287516076	Test Loss: 0.7980399131343402
Epoch:42	Train Loss:0.5535261519360201	Test Loss: 0.7760960974201963
Epoch:43	Train Loss:0.5572962617831367	Test Loss: 0.7732255493999103
Epoch:44	Train Loss:0.5482994141725547	Test Loss: 0.7805390681452963
Epoch:45	Train Loss:0.5496758268245955	Test Loss: 0.7748767138246634
Epoch:46	Train Loss:0.5449914050444461	Test Loss: 0.7937413368906293
Epoch:47	Train Loss:0.5455981116082395	Test Loss: 0.7741713747094496
Epoch:48	Train Loss:0.5426202269618591	Test Loss: 0.7653982825577259
Epoch:49	Train Loss:0.5415597493033946	Test Loss: 0.810625221939851
Epoch:50	Train Loss:0.5413240872822594	Test Loss: 0.7864760395973691

```
plt.plot(range(len(epoch_losses)), epoch_losses, c="orange",  
label="train loss")  
plt.plot(range(len(test_epoch_losses)), test_epoch_losses, c="blue",  
label="test loss")  
plt.legend("top right")  
plt.show()
```



```

results_final = pd.DataFrame(columns=['actual', 'prediction', 'dset'])

with torch.no_grad():
    for image_batch, text_batch in tqdm(train_loader, leave=True):
        text_batch_logits = crnn(image_batch.to(device)) # [T,
        batch_size, num_classes==num_features]
        text_batch_pred = decode_predictions(text_batch_logits.cpu())
        #print(text_batch, text_batch_pred)
        df = pd.DataFrame(columns=['actual', 'prediction'])
        df['actual'] = text_batch
        df['prediction'] = text_batch_pred
        df['dset'] = "train"
        results_final = pd.concat([results_final, df])
results_final = results_final.reset_index(drop=True)

{"model_id": "8849fefe095d404b88ae84db56cdc0f4", "version_major": 2, "version_minor": 0}

with torch.no_grad():
    for image_batch, text_batch in tqdm(test_loader, leave=True):
        text_batch_logits = crnn(image_batch.to(device)) # [T,
        batch_size, num_classes==num_features]
        text_batch_pred = decode_predictions(text_batch_logits.cpu())
        #print(text_batch, text_batch_pred)
        df = pd.DataFrame(columns=['actual', 'prediction'])
        df['actual'] = text_batch

```



```

df['prediction'] = text_batch_pred
df['dset'] = "test"
results_final = pd.concat([results_final, df])
results_final = results_final.reset_index(drop=True)

{"model_id": "27b290896ef048b3bd00fe009c570117", "version_major": 2, "version_minor": 0}

results_final['prediction_corrected'] =
results_final['prediction'].apply(correct_prediction)
results_final

```

	actual	prediction	dset	prediction_corrected
0	forced	foo-r-c-e-d--	train	forced
1	the	t---h----e---	train	the
2	den	d----e--n----	train	den
3	the	t---h----e---	train	the
4	uit	u-----i--t--	train	uit
...
23482	de	d-----e-----	test	de
23483	heeft	h---e-e--f-t-	test	heeft
23484	in	i----n-----	test	in
23485	dat	d----a----t--	test	dat
23486	ookats	o-o-k--a--l--	test	ookal

[23487 rows x 4 columns]

```

import json
from weighted_levenshtein import lev
import numpy as np

with open("./params_weighted_leven.json", "r") as f:
    leven_params = json.load(f)
    for k in leven_params.keys():
        leven_params[k] = np.array(leven_params[k])

results_final["val"] = results_final.apply(lambda row:
lev(row["actual"], row["prediction_corrected"], **leven_params),
axis=1)

results_final

```

	actual	prediction	dset	prediction_corrected	val
0	forced	foo-r-c-e-d--	train	forced	0.0000
1	the	t---h----e---	train	the	0.0000
2	den	d----e--n----	train	den	0.0000
3	the	t---h----e---	train	the	0.0000
4	uit	u-----i--t--	train	uit	0.0000
...
23482	de	d-----e-----	test	de	0.0000
23483	heeft	h---e-e--f-t-	test	heeft	0.0000

23484	in	i-----	test	in	0.0000
23485	dat	d----a----t--	test	dat	0.0000
23486	ookats	o-o-k--a--l--	test	ookal	1.7252

[23487 rows x 5 columns]